

# Linear, Multiple, and Logistic Regressions

CMPUT 328

Nilanjan Ray

Computing Science, University of Alberta, Canada

# Linear regression with PyTorch

- We will start with a linear regression “model”
- Next, we need to understand “loss” function for regression task
- Next, we will estimate the model by minimizing the loss function
- We will use PyTorch

# Supervised machine learning: the tabular view

Independent variables (aka <b>feature vector</b> )				Prediction / dependent variable
$x_1$	$x_2$	$x_3$	$x_4$	$y$
1.2	-3.9	4.0	0	1.6
2.1	2.4	-0.7	-0.2	1.2
...	...	...	...	...
...	...	...	...	...
3.2	...	...	1.9	0.3
1.4	...	...	1.5	?
3.1	...	...	2.1	?



ML learns to map **x to y**

Training data:  
**complete** table

In other words, ML learns  
a function,  $f$  so that  
 $y = f(x)$

Test data:  
**incomplete** table

The function  $f$  is called **prediction function**

# Linear prediction: formal setup

Linear prediction function:

$$y^p = \mathbf{x}\boldsymbol{\theta} + b$$

or,

$$y^p = \sum_{j=1}^m \theta_j x_j + b$$

vector equation form

scalar equation form

A training set consists of  $(\mathbf{x}, y)$  pairs:  $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$

Linear prediction on the training data point  $i$ :

$$y_i^p = \mathbf{x}_i \boldsymbol{\theta} + b$$

or,

$$y_i^p = \sum_{j=1}^m \theta_j x_{i,j} + b$$

Loss or cost function (on training data):

$$L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 = \frac{1}{2} \sum_{i=1}^n (\mathbf{x}_i \boldsymbol{\theta} + b - y_i)^2$$

Learning the linear model: find out  $\boldsymbol{\theta}$  and  $b$  to minimize loss function

# Linear regression: A toy example

Let's take a toy example:

$x_1$	$x_2$	$y$
1	2	-1
3	-4	7
6	2	3
-3	5	-4
7	-3	5
4	3	?

This equation  $y_i^p = \sum_{j=1}^m \theta_j x_{i,j} + b$

can be written for the toy training set as

We also have ground truth responses:

$$y_1 = -1, y_2 = 7, y_3 = 3, y_4 = -4, y_5 = 5$$

$$\begin{aligned}y_1^p &= \theta_1(1) + \theta_2(2) + b \\y_2^p &= \theta_1(3) + \theta_2(-4) + b \\y_3^p &= \theta_1(6) + \theta_2(2) + b \\y_4^p &= \theta_1(-3) + \theta_2(5) + b \\y_5^p &= \theta_1(7) + \theta_2(-3) + b\end{aligned}$$

So, the loss is  $L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 = \frac{1}{2} [(y_1^p + 1)^2 + (y_2^p - 7)^2 + (y_3^p - 3)^2 + (y_4^p + 4)^2 + (y_5^p - 5)^2]$

$$= \frac{1}{2} [(\theta_1 + 2\theta_2 + b + 1)^2 + (3\theta_1 - 4\theta_2 + b - 7)^2 + (6\theta_1 + 2\theta_2 + b - 3)^2 + (-3\theta_1 + 5\theta_2 + b + 4)^2 + (7\theta_1 - 3\theta_2 + b - 5)^2]$$

# Learning a linear model

For the convenience of math, let us change our linear model a bit:

$$y_i^p = \sum_{j=1}^m \theta_j(x_{i,j} - \bar{x}_j) + b \quad \text{where} \quad \bar{x}_j = \frac{1}{n} \sum_{i=1}^n x_{i,j}$$

Subtracting independent variable mean is called “data centering”

Why? Because it simplifies the math later when solving for  $\theta$  and  $b$ .

And a slightly modified loss function:

$$L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 + \frac{\gamma}{2} \sum_{j=1}^m \theta_j^2$$

$\gamma$  is a **hyper parameter**



Data fidelity

Regularization

Why do we need regularization?

# Minimization of linear regression loss function

Regularized loss function:  $L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 + \frac{\gamma}{2} \sum_{j=1}^m \theta_j^2$

Taking partial derivative using chain rule:  $\frac{\partial L}{\partial b} = \sum_{i=1}^n (y_i^p - y_i) \frac{\partial y_i^p}{\partial b} = \sum_{i=1}^n (y_i^p - y_i)$  because,  $\frac{\partial y_i^p}{\partial b} = 1$

Using  $\bar{x}_j = \frac{1}{n} \sum_{i=1}^n x_{i,j}$  and  $y_i^p = \sum_{j=1}^m \theta_j (x_{i,j} - \bar{x}_j) + b$  we get:  $\frac{\partial L}{\partial b} = nb - \sum_{i=1}^n y_i$

At the minimum of  $L$ ,  $\frac{\partial L}{\partial b} = 0$  So,  $b = \frac{1}{n} \sum_{i=1}^n y_i = \bar{y}$

# Linear regression: A toy example...continued

Let's take a toy example:

$x_1$	$x_2$	y
1	2	-1
3	-4	7
6	2	3
-3	5	-4
7	-3	5
4	3	?

$$b = \frac{1}{n} \sum_{i=1}^n y_i = \bar{y} = \frac{1}{5}(-1 + 7 + 3 - 4 + 5) = 2$$

$$\bar{x}_1 = \frac{1}{n} \sum_{i=1}^n x_{i,1} = \frac{1}{5}(1 + 3 + 6 - 3 + 7) = 2.8$$

$$\bar{x}_2 = \frac{1}{n} \sum_{i=1}^n x_{i,2} = \frac{1}{5}(2 - 4 + 2 + 5 - 3) = 0.4$$

So, using centered data, the prediction equation becomes:

$$y_i^p = \sum_{j=1}^m \theta_j (x_{i,j} - \bar{x}_j) + b = \theta_1(x_{i,1} - 2.8) + \theta_2(x_{i,2} - 0.4) + 2$$

So, the loss is

$$\begin{aligned}
 L &= \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 = \frac{1}{2} [(y_1^p + 1)^2 + (y_2^p - 7)^2 + (y_3^p - 3)^2 + (y_4^p + 4)^2 + (y_5^p - 5)^2] \\
 &= \frac{1}{2} [(\theta_1(1 - 2.8) + \theta_2(2 - 0.4) + 2 + 1)^2 + (\theta_1(3 - 2.8) + \theta_2(-4 - 0.4) + 2 - 7)^2 \\
 &\quad + (\theta_1(6 - 2.8) + \theta_2(2 - 0.4) + 2 - 3)^2 + (\theta_1(-3 - 2.8) + \theta_2(5 - 0.4) + 2 + 4)^2 + (\theta_1(7 - 2.8) + \theta_2(-3 - 0.4) + 2 - 5)^2]
 \end{aligned}$$

# Minimization of linear regression loss function...

Regularized loss function:

$$L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 + \frac{\gamma}{2} \sum_{j=1}^m \theta_j^2$$

Taking partial derivative of  $L$  using chain rule:

$$\frac{\partial L}{\partial \theta_j} = \sum_{i=1}^n (y_i^p - y_i) \frac{\partial y_i^p}{\partial \theta_j} + \gamma \theta_j$$

Using

$$y_i^p = \sum_{k=1}^m \theta_k (x_{i,k} - \bar{x}_k) + b,$$

$$b = \bar{y}$$

and

$$\frac{\partial y_i^p}{\partial \theta_j} = x_{i,j} - \bar{x}_j$$

We get:  $\frac{\partial L}{\partial \theta_j} = \sum_{i=1}^n \left( \sum_{k=1}^m \theta_k (x_{i,k} - \bar{x}_k) + \bar{y} - y_i \right) (x_{i,j} - \bar{x}_j) + \gamma \theta_j$

# Linear regression: A toy example...continued

Let's take a toy example:

$x_1$	$x_2$	$y$
1	2	-1
3	-4	7
6	2	3
-3	5	-4
7	-3	5
4	3	?

Note: For this toy problem, I assumed  $\gamma = 0$  for convenience

$$\frac{\partial L}{\partial \theta_j} = \sum_{i=1}^n \left( \sum_{k=1}^m \theta_k (x_{i,k} - \bar{x}_k) + \bar{y} - y_i \right) (x_{i,j} - \bar{x}_j) + \gamma \theta_j$$

$$\begin{aligned}\frac{\partial L}{\partial \theta_1} &= (\theta_1(1 - 2.8) + \theta_2(2 - 0.4) + 2 + 1)(1 - 2.8) \\ &+ (\theta_1(3 - 2.8) + \theta_2(-4 - 0.4) + 2 - 7)(3 - 2.8) \\ &+ (\theta_1(6 - 2.8) + \theta_2(2 - 0.4) + 2 - 3)(6 - 2.8) \\ &+ (\theta_1(-3 - 2.8) + \theta_2(5 - 0.4) + 2 + 4)(-3 - 2.8) \\ &+ (\theta_1(7 - 2.8) + \theta_2(-3 - 0.4) + 2 - 5)(7 - 2.8) \\ &= (64.8)\theta_1 - (39.6)\theta_2 - 57\end{aligned}$$

$$\begin{aligned}\frac{\partial L}{\partial \theta_2} &= (\theta_1(1 - 2.8) + \theta_2(2 - 0.4) + 2 + 1)(2 - 0.4) \\ &+ (\theta_1(3 - 2.8) + \theta_2(-4 - 0.4) + 2 - 7)(-4 - 0.4) \\ &+ (\theta_1(6 - 2.8) + \theta_2(2 - 0.4) + 2 - 3)(2 - 0.4) \\ &+ (\theta_1(-3 - 2.8) + \theta_2(5 - 0.4) + 2 + 4)(5 - 0.4) \\ &+ (\theta_1(7 - 2.8) + \theta_2(-3 - 0.4) + 2 - 5)(-3 - 0.4) \\ &= (-39.6)\theta_1 + (57.2)\theta_2 + 63\end{aligned}$$

# Minimization of linear regression loss function...

$$\frac{\partial L}{\partial \theta_j} = \sum_{i=1}^n \left( \sum_{k=1}^m \theta_k (x_{i,k} - \bar{x}_k) + \bar{y} - y_i \right) (x_{i,j} - \bar{x}_j) + \gamma \theta_j$$

↓  
simplification

Gradient of  $L$  with  
resp. to  $\boldsymbol{\theta}$ :

$$\nabla_{\boldsymbol{\theta}} L = \left[ \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}})^T (\mathbf{x}_i - \bar{\mathbf{x}}) \right] \boldsymbol{\theta} + \gamma \boldsymbol{\theta} - \sum_{i=1}^n (y_i - \bar{y}) (\mathbf{x}_i - \bar{\mathbf{x}})^T$$

where  $\mathbf{x}_i = [x_{i,1} \dots x_{i,m}]$ ,  $\bar{\mathbf{x}} = [\bar{x}_1 \dots \bar{x}_m]$  and  $\boldsymbol{\theta} = [\theta_1 \dots \theta_m]^T$

More simplified form:  $\nabla_{\boldsymbol{\theta}} L = (X^T X + \gamma I) \boldsymbol{\theta} - X^T \mathbf{y}$

where matrix  $X$  is defined as:  $X = \begin{bmatrix} \mathbf{x}_1 - \bar{\mathbf{x}} \\ \vdots \\ \mathbf{x}_n - \bar{\mathbf{x}} \end{bmatrix}$  and vector  $\mathbf{y}$  is defined as:  $\mathbf{y} = \begin{bmatrix} y_1 - \bar{y} \\ \vdots \\ y_n - \bar{y} \end{bmatrix}$

and  $I$  is an identity matrix of size  $m$ -by- $m$

Equating gradient of  $L$  to zero vector and solving for  $\boldsymbol{\theta}$  gives us:

$$\boldsymbol{\theta} = (X^T X + \gamma I)^{-1} X^T \mathbf{y}$$

# Quick review: Gradient of a function

Consider a function of two variables as an example:

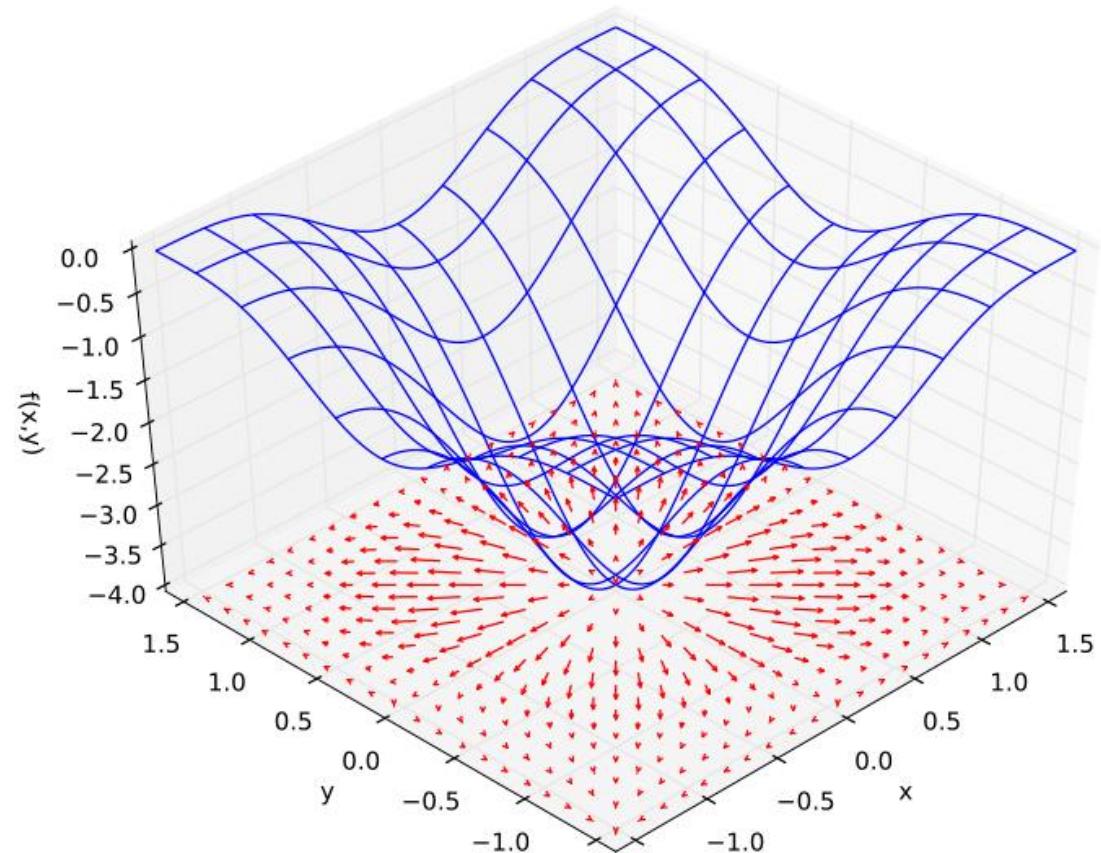
$$f(x, y) = -(\cos^2 x + \cos^2 y)^2$$

$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix} = \begin{bmatrix} 4(\cos^2(x) + \cos^2(y)) \cos(x) \sin(x) \\ 4(\cos^2(x) + \cos^2(y)) \cos(y) \sin(y) \end{bmatrix}$$

**Note 1:**  $f$  is a function of **two variables**,  
so gradient (partial derivatives collected as a vector)  
of  $f$  is a **two-dimensional vector**

**Note 2:** Gradient (vector) of  $f$  points toward the  
**steepest ascent for  $f$**

**Note 3:** At a (local) minimum of  $f$  its gradient  
becomes a **zero vector**



Example source: Wikipedia

# Gradient in our toy example...

Toy example:

$x_1$	$x_2$	y
1	2	-1
3	-4	7
6	2	3
-3	5	-4
7	-3	5
4	3	?

Note: For this toy problem, I assumed  $\gamma = 0$  for convenience

$$\frac{\partial L}{\partial \theta_1} = (64.8)\theta_1 - (39.6)\theta_2 - 57$$

$$\frac{\partial L}{\partial \theta_2} = (-39.6)\theta_1 + (57.2)\theta_2 + 63$$

So, gradient is

$$\nabla L_{\theta} = \begin{bmatrix} (64.8)\theta_1 - (39.6)\theta_2 - 57 \\ (-39.6)\theta_1 + (57.2)\theta_2 + 63 \end{bmatrix}$$

Equating gradient (vector) to 0 (vector) and solving, we get:

$$\theta = \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix} = \begin{bmatrix} 0.3580 \\ -0.8535 \end{bmatrix}$$

# Equivalently, using direct formula...

Let's take a toy example:

$x_1$	$x_2$	y
1	2	-1
3	-4	7
6	2	3
-3	5	-4
7	-3	5
4	3	?

$$X = \begin{bmatrix} \mathbf{x}_1 - \bar{\mathbf{x}} \\ \vdots \\ \mathbf{x}_n - \bar{\mathbf{x}} \end{bmatrix} = \begin{bmatrix} 1 - 2.8 & 2 - 0.4 \\ 3 - 2.8 & -4 - 0.4 \\ 6 - 2.8 & 2 - 0.4 \\ -3 - 2.8 & 5 - 0.4 \\ 7 - 2.8 & -3 - 0.4 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 - \bar{y} \\ \vdots \\ y_n - \bar{y} \end{bmatrix} = \begin{bmatrix} -1 - 2 \\ 7 - 2 \\ 3 - 2 \\ -4 - 2 \\ 5 - 2 \end{bmatrix}$$

$$\boldsymbol{\theta} = (X^T X)^{-1} X^T \mathbf{y} = \begin{bmatrix} 0.3580 \\ -0.8535 \end{bmatrix}$$

So, finally the prediction for the test data point

$$? = \sum_{j=1}^m \theta_j (x_j - \bar{x}_j) + b = 0.3580(4 - 2.8) - 0.8535(3 - 0.4) + 2 = 0.2105$$

Note: For this problem I did not add any regularization, i.e., I assumed  $\gamma = 0$

# MNIST Dataset



Classify images into digits

Each image is **28x28**

**10** labels

**55,000** training images

**5,000** validation images

**10,000** test images.

# Linear regression on MNIST dataset



Small 28 pixels-by-28 pixels images of hand written digits

The visual recognition problem definition:  
**to recognize the digit from an image**

Our very first line of attack would be to  
use linear regression.

Feature dimension,  $m = 28 * 28 = 784$

Let's look at our PyTorch implementations:  
Called direct method because we will use math formula  
to find  $\theta$  and b.

Pixel values (feature)					Digit
$x_1$	$x_2$	...	$x_{784}$	$y$	
0.1	0.3	...	0.0	0	
0.2	0.1	...	0.5	1	
...	...	...	...	...	
...	...	...	...	...	
0.0	0.98	...	0.8	9	
0.5	0.25	...	0.36	?	
0.1	0.95	...	0.1	?	

# Linear regression

So, far we have seen:



[28x28]

Array of real numbers  
(784 numbers in total)

Image

(1x784)

Parameters

(784x1)

$$\mathbf{y}^p = (\mathbf{x} - \bar{\mathbf{x}}) \boldsymbol{\theta} + \bar{y}$$

Mean vector  
of training  
images

Mean of  
training  
labels (digits)

1 number,  
indicating digit

See notebook: MNIST\_Linear\_Regression\_Direct.ipynb

Pixel values (feature)				Digit
$x_1$	$x_2$	...	$x_{784}$	$y$
0.1	0.3	...	0.0	0
0.2	0.1	...	0.5	1
...	...	...	...	...
...	...	...	...	...
0.0	0.98	...	0.8	9
0.5	0.25	...	0.36	?
0.1	0.95	...	0.1	?

# Multiple or Vector Linear Regression



**[28x28]**  
Array of real  
numbers (784  
numbers total)

Image  
(1x784)

Parameters  
(784x10)

$$\mathbf{y}^p = (\mathbf{x} - \bar{\mathbf{x}}) \mathbf{W} + \bar{\mathbf{y}}$$

Mean  
vector  
of  
training  
images

Mean of  
1-hot  
training  
label  
vector

10 numbers,  
indicating class  
scores

Pixel values (feature)      Digit: 1-hot vector

$x_1$	$x_2$	...	$x_{784}$	$y_1$	...	$y_{10}$
0.1	0.3	...	0.0	0	...	1
0.2	0.1	...	0.5	1		0
...	...	...	...	...	...	...
...	...	...	...	...	...	...
0.0	0.98	...	0.8	0	...	1
0.5	0.25	...	0.36	?	...	?
0.1	0.75	...	0.1	?	...	?

# Multiple Linear Regression: PyTorch Implementation

See notebook: MNIST\_Multiple\_Linear\_Regression\_Direct.ipynb

Prediction model:  $\mathbf{y}^p = (\mathbf{x} - \bar{\mathbf{x}})W + \bar{\mathbf{y}}$

[https://en.wikipedia.org/wiki/Matrix\\_calculus](https://en.wikipedia.org/wiki/Matrix_calculus)

Regularized loss function:  $L = \frac{1}{2} \sum_{i=1}^n \|\mathbf{y}_i^p - \mathbf{y}_i\|^2 + \frac{\gamma}{2} \|W\|^2$

Gradient of loss function:

$$\nabla_W L = (X^T X + \gamma I)W - X^T Y$$

where matrix  $X$  is defined as:

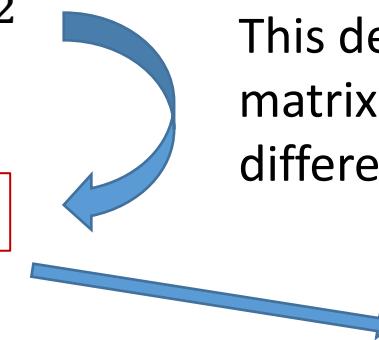
$$X = \begin{bmatrix} \mathbf{x}_1 - \bar{\mathbf{x}} \\ \vdots \\ \mathbf{x}_n - \bar{\mathbf{x}} \end{bmatrix}$$

and matrix  $Y$  is defined as:  $Y = \begin{bmatrix} \mathbf{y}_1 - \bar{\mathbf{y}} \\ \vdots \\ \mathbf{y}_n - \bar{\mathbf{y}} \end{bmatrix}$

and  $I$  is an identity matrix of size 784-by-784

We will “minimally” modify our linear regression scripts into multiple linear regression implementations!

This derivation requires matrix-vector differentiation



Equating gradient of  $L$  to zero matrix and solving for  $W$  gives us:

$$W = (X^T X + \gamma I)^{-1} X^T Y$$

# What happened to our learning algorithm?

- Step 1: Create training image set (example set):

```
0000000000000000  
1111111111111111  
2222222222222222  
3333333333333333  
4444444444444444  
5555555555555555  
6666666666666666  
7777777777777777  
8888888888888888  
9999999999999999
```

Repeat steps 2, 3 and 4

- Step 2: Show these examples to the machine learner
- Step 3: Measure mistakes made by the machine learner
- Step 4: Tune parameters of the machine learner to minimize its mistakes

# Can we apply this learning algorithm to linear or multiple linear regression using PyTorch?

Iterate:

(Load Data): Get a training data batch (also called mini batch)

(Predict): Apply linear model to training feature vector and compute predictions

(Compute loss): Measure discrepancy between predictions and ground truths

(Optimize): Ask PyTorch to reduce loss value by tuning the parameters  $\theta$  (or  $W$ ), and  $b$

(Diagnostics): Check if loss is decreasing

See notebook: MNIST\_Linear\_Regression.ipynb, MNIST\_Multiple\_Linear\_Regression.ipynb

What are the pros and cons of this optimization-based method over the direct formula-based method?

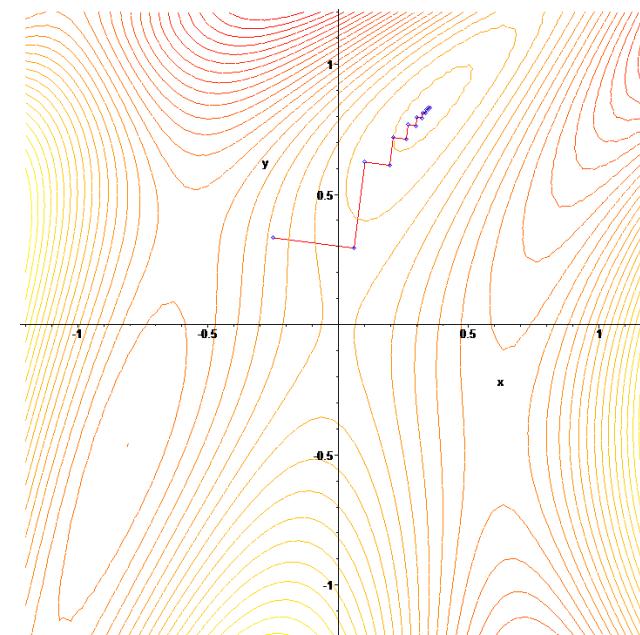
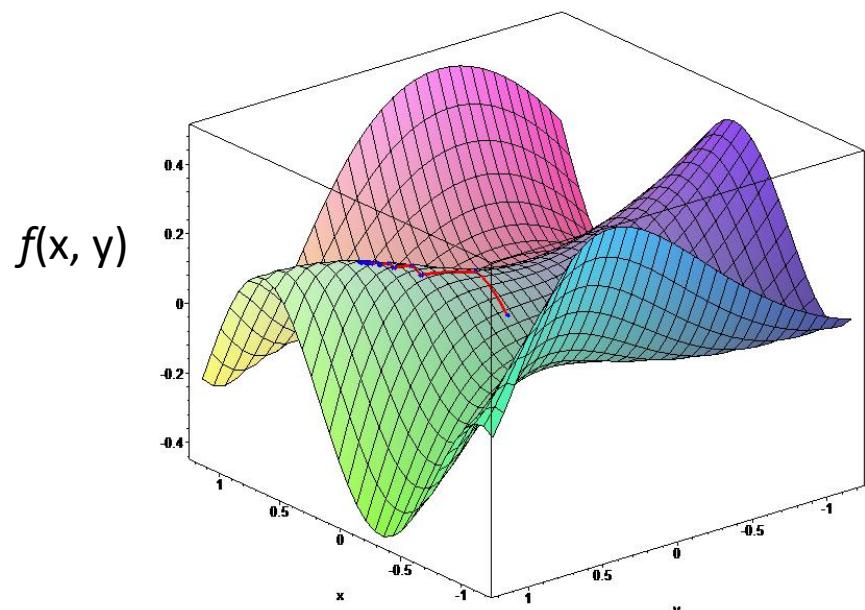
# Gradient descent optimization

Start at an initial guess for the optimization variable:  $\mathbf{x}_0$

Iterate until gradient magnitude becomes too small:  $\mathbf{x}^{t+1} = \mathbf{x}^t - \alpha \nabla f(\mathbf{x}^t)$

} Gradient descent algorithm

$\alpha$  is called the step-length.



Gradient descent creates a zig-zag path leading to a local minimum of  $f$

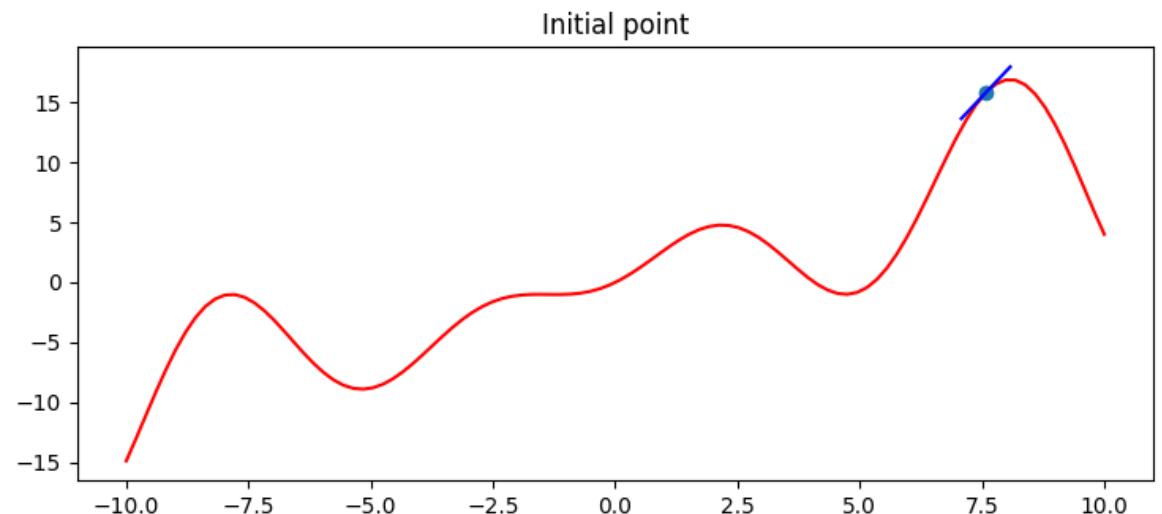
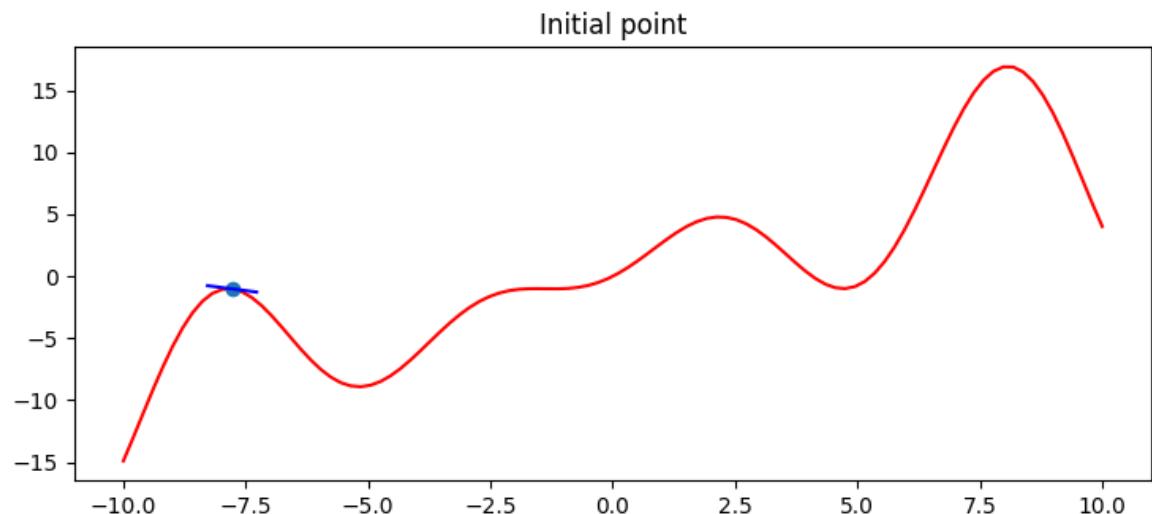
# Gradient descent visualization

Find  $x$  such that  $f(x)$  is minimized:  $f(x) = \sin(x) + x + x * \sin(x)$

$$\nabla f(x) = \cos(x) + 1 + \sin(x) + x * \cos(x)$$

$$\mathbf{x}_0 = ???$$

$$\mathbf{x}^{t+1} = \mathbf{x}^t - \alpha \nabla f(\mathbf{x}^t)$$



# PyTorch optimizer uses GD

Let's try our own gradient descent for multiple linear regression

Gradient of loss function for multiple linear regression:  $\nabla_W L = (X^T X + \gamma I)W - X^T Y$

$$\frac{\partial L}{\partial b} = \sum_{i=1}^n (y_i^p - y_i)$$

Exercise: write GD for MNIST multiple linear regression

Look at `MNIST_Multiple_Linear_Regression.ipynb`

# Logistic Regression

Can we modify scores from multiple regression function to output probabilities?

What is a suitable loss function for classification?

# Logistic Regression

Would it not be nice if we can predict **class probabilities** instead of scores?



[28x28]

Array of real numbers  
(784 numbers total)

image      parameters

$$\mathbf{y}^p = f(\mathbf{x}, \mathbf{W})$$

prediction function  
For logistic regression

10 numbers,  
indicating class  
probabilities

Pixel values (feature)						Digit: 1-hot vector
$x_1$	$x_2$	...	$x_{784}$	$y_1$	...	$y_{10}$
0.1	0.3	...	0.0	0	...	1
0.2	0.1	...	0.5	1	...	0
...	...	...	...	...	...	...
...	...	...	...	...	...	...
0.0	0.98	...	0.8	0	...	1
0.5	0.25	...	0.36	?	...	?
0.4	0.95	...	0.1	?	...	?

# Logistic regression: from multiple linear regression

Scores from multiple linear regression:

$$\mathbf{s}_i = (\mathbf{x}_i - \bar{\mathbf{x}})W + \bar{\mathbf{y}}$$

$$\mathbf{s}_{i,k} = (\mathbf{x}_i - \bar{\mathbf{x}})W_{:,k} + \bar{\mathbf{y}}_k$$

Score for  $k^{\text{th}}$  class,  $k = 0, \dots, 9$

Predicted probability for  $k^{\text{th}}$  class:

$$y_{i,k}^p = \frac{\exp(s_{i,k})}{\sum_{c=0}^9 \exp(s_{i,c})}$$

“Softmax” function

# Logistic regression: loss function

Cross entropy loss: 
$$loss(\mathbf{y}^p, \mathbf{y}) = - \sum_{k=0}^9 \mathbf{y}_k \log(\mathbf{y}_k^p)$$

Why this loss function? What does it mean? Why not use Euclidean loss as in MLR?

Do we have a direct formula to compute parameters like MLR?